

Probability, Likelihood, and Other Measures of Model Fit


Mark Andrews

Psychology Department, Nottingham Trent University

✉ mark.andrews@ntu.ac.uk

Example problem

- Let's assume we have the cars data, which is depicted in the following scatterplot:



Example problem

- The first 10 observations of cars are:

```
head(cars, 10)
#>      speed dist
#> 1         4    2
#> 2         4   10
#> 3         7    4
#> 4         7   22
#> 5         8   16
#> 6         9   10
#> 7        10   18
#> 8        10   26
#> 9        10   34
#> 10       11   17
```

Probabilistic model

- ▶ A potential model of the cars data is the following

$$y_i \sim N(\mu_i, \sigma^2) \quad \text{for } i \in 1 \dots n,$$
$$\mu_i = \beta_0 + \beta_1 x_i,$$

where y_i and x_i are the dist and speed variables on observation i .

- ▶ In other words, we are modelling dist as normally distributed around a mean that is a linear function of speed, and with a fixed variance σ^2 .
- ▶ We do not know the values of the parameters β_0 , β_1 , and σ^2 .
- ▶ Note that this is a probabilistic model of the outcome variable only.

Conditional probability of any observation

- ▶ Given our model specification, we can ask what is the probability of any given value of `dist`, assuming a given value of `speed`, for any given values of the parameters $\beta_0, \beta_1, \sigma^2$.
- ▶ For example, we can ask, what is the probability that `dist` = 50 if `speed` = 15 if β_0, β_1 , and σ have the values $-20, 4, 15$, respectively, i.e.

$$P(y = 50 | x = 15, \beta_0 = -20, \beta_1 = 4, \sigma = 15),$$

- ▶ We can do this for *any* values of `dist`, `speed`, and β_0, β_1 , and σ .

Conditional probability of any observation

- If $x = 15$, and $\beta_0 = -20$, $\beta_1 = 4$, $\sigma = 15$, then the value of y has been assumed to be drawn from a normal distribution with mean

$$\mu = \beta_0 + \beta_1 x,$$

$$\mu = -20 + 4 \times 15,$$

$$\mu = 40$$

and a standard deviation of $\sigma = 15$.

- And so the probability that $y = 50$ when $x = 15$, and $\beta_0 = -20$, $\beta_1 = 4$, $\sigma = 15$, is the probability of a value of 50 in a normally distributed random variable whose mean is 40 and whose standard deviation is 15.

Conditional probability of any observation

- The probability (density) that a normal random variable, with mean of 40 and standard deviation of 15, takes the value of 50 can be obtained from this probability density function for normal distributions:

$$P(y|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{|y - \mu|^2}{2\sigma^2}\right)$$

- Using R, this is

```
y <- 50; mu <- 40; sigma <- 15  
1/sqrt(2*pi*sigma^2) * exp(-0.5 * (y-mu)^2/sigma^2)  
#> [1] 0.02129653
```

or just

```
dnorm(y, mean = mu, sd = sigma)  
#> [1] 0.02129653
```

Conditional probability of any observation

- We can use the R function `prob_obs_lm` (from `utils.R`) for the probability of an observation in any (simple) linear regression:

```
# e.g.  
prob_obs_lm(y = 50,  
            x = 15,  
            beta_0 = -20, beta_1 = 4, sigma = 15)  
#> [1] 0.02129653
```


Conditional probability of all observed data

- ▶ Assuming values for β_0, β_1, σ , what the probability of the observed values of the dist outcome variable, $y_1, y_2, y_3 \dots y_n$ given the observed values of the speed predictor, $x_1, x_2, x_3 \dots x_n$?
- ▶ This is

$$P(y_1 \dots y_n | x_1 \dots x_n, \beta_0, \beta_1, \sigma).$$

- ▶ In this model, all y 's are conditionally independent of one another, given that values of $x_1, x_2, x_3 \dots x_n$, so the the joint probability is as follows:

$$P(y_1 \dots y_n | x_1 \dots x_n, \beta_0, \beta_1, \sigma) = \prod_{i=1}^n P(y_i | x_i, \beta_0, \beta_1, \sigma).$$

Conditional log probability of all observed data

- The joint probability

$$P(y_1 \dots y_n | x_1 \dots x_n, \beta_0, \beta_1, \sigma) = \prod_{i=1}^n P(y_i | x_i, \beta_0, \beta_1, \sigma).$$

will be a very small number (a product of small numbers), so we usually calculate its logarithm:

$$\log \left(\prod_{i=1}^n P(y_i | x_i, \beta_0, \beta_1, \sigma) \right) = \sum_{i=1}^n \log P(y_i | x_i, \beta_0, \beta_1, \sigma)$$

Log conditional probability of all observed data

- For example, the log probability of all the dist values given the speed values, and assuming certain values for β_0 , β_1 , σ can be calculated as follows:

```
y <- cars$dist; x <- cars$speed
beta_0 = -20; beta_1 = 4; sigma = 15
prob_obs_lm(y, x, beta_0, beta_1, sigma, log = TRUE) %>%
  sum()
#> [1] -206.805
```

Log conditional probability of all observed data

- The log conditional probability can

$$\begin{aligned}\sum_{i=1}^n \log P(y_i | x_i, \beta_0, \beta_1, \sigma) &= \sum_{i=1}^n \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{|y_i - \mu_i|^2}{2\sigma^2} \right) \right), \\ &= -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n |y_i - \mu_i|^2,\end{aligned}$$

where $\mu_i = \beta_0 + \beta_1 x_i$.

- This is calculated by `log_prob_obs_lm` (in `utils.R`):

```
log_prob_obs_lm(y, x, beta_0, beta_1, sigma)
#> [1] -206.805
```

The likelihood function

- The following is a function over the space of values of $y_1 \dots y_n$:

$$P(y_1 \dots y_n | x_1 \dots x_n, \beta_0, \beta_1, \sigma) = \prod_{i=1}^n P(y_i | x_i, \beta_0, \beta_1, \sigma).$$

- In other words, $x_1 \dots x_n$ and β_0, β_1 , and σ are fixed (like the parameters of a function) and $y_1 \dots y_n$ are free variables and so $P(y_1 \dots y_n | x_1 \dots x_n, \beta_0, \beta_1, \sigma)$ is a function over the $y_1 \dots y_n$ space.
- If, however, we treat $y_1 \dots y_n$ and $x_1 \dots x_n$ as fixed, and treat β_0, β_1 , and σ as free variables, then

$$\mathcal{L}(\beta_0, \beta_1, \sigma | \vec{y}, \vec{x}) = \prod_{i=1}^n P(y_i | x_i, \beta_0, \beta_1, \sigma)$$

defines a function over the three dimensional β_0, β_1, σ space.

- The function is known as the *likelihood function*.

The log likelihood function

- ▶ The log likelihood function is just the log of the likelihood function.
- ▶ In the present example, it is

$$\begin{aligned} \log \mathcal{L}(\beta_0, \beta_1, \sigma | \vec{y}, \vec{x}) = \\ -\frac{n}{2} \log(2\pi) - \frac{n}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n |y_i - (\beta_0 + \beta_1 x_i)|^2, \end{aligned}$$

where $y_1 \dots y_n$ and $x_1 \dots x_n$ are assumed to be fixed.

Maximum likelihood estimation

- ▶ The values of β_0, β_1, σ that maximize $\mathcal{L}(\beta_0, \beta_1, \sigma | \vec{y}, \vec{x})$ are the maximum likelihood estimators of the (random variables) β_0, β_1, σ .
- ▶ The values of β_0, β_1, σ that maximize $\log \mathcal{L}(\beta_0, \beta_1, \sigma | \vec{y}, \vec{x})$ are those that maximize $\mathcal{L}(\beta_0, \beta_1, \sigma | \vec{y}, \vec{x})$.
- ▶ We usually denote estimators by $\hat{\beta}_0, \hat{\beta}_1, \hat{\sigma}$.
- ▶ By definition, the maximum likelihood estimators $\hat{\beta}_0, \hat{\beta}_1, \hat{\sigma}$ are the values of β_0, β_1, σ that maximize the probability of the observed data.

Maximum likelihood estimation

- ▶ In a simple linear regression, the maximum likelihood estimators for the linear coefficients are those that minimize the following

$$\sum_{i=1}^n |y_i - (\beta_0 + \beta_1 x_i)|^2,$$

- ▶ In general, for linear regression, the maximum likelihood estimators always minimize the sum of squared residuals.
- ▶ In R, for the cars data, the maximum likelihood estimators for β_0 and β_1 are obtained as follows:

```
M1 <- lm(dist ~ speed, data = cars)
coef(M1)
#> (Intercept)      speed
#>  -17.579095    3.932409
```


Maximum likelihood estimation

- ▶ The maximum likelihood estimator for σ^2 is the following:

$$\frac{1}{n} \sum_{i=1}^n |y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)|^2,$$

- ▶ Using R, we can obtain this as follows:

```
mean(residuals(M1)^2)
#> [1] 227.0704
```

and so the maximum likelihood estimator for σ is

```
mean(residuals(M1)^2) %>% sqrt()
#> [1] 15.06886
```

The model's log likelihood

- ▶ When we speak of the log likelihood of a model, we mean the maximum value of the model's log likelihood function.
- ▶ In other words, it is the value of the log likelihood function at its maximum likelihood estimators' values.
- ▶ In yet other words, **it is the (log) probability of the observed data given the maximum likelihood estimates of its parameters.**

The model's log likelihood

```
beta_0_mle <- coef(M1)['(Intercept)']
beta_1_mle <- coef(M1)['speed']
sigma_mle <- mean(residuals(M1)^2) %>% sqrt()

log_prob_obs_lm(y, x,
                beta_0 = beta_0_mle,
                beta_1 = beta_1_mle,
                sigma = sigma_mle) %>% sum()

#> [1] -206.5784
# same as ...
logLik(M1)
#> 'log Lik.' -206.5784 (df=3)
```

Residual sum of squares

- ▶ The sum of squared residuals in a simple linear model is

$$\text{RSS} = \sum_{i=1}^n |y_i - (\beta_0 + \beta_1 x_i)|^2.$$

- ▶ The RSS when using the maximum likelihood estimators is

$$\begin{aligned}\text{RSS} &= \sum_{i=1}^n |y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)|^2, \\ &= \sum_{i=1}^n |y_i - \hat{y}_i|^2\end{aligned}$$

Residual sum of squares and log likelihood

- ▶ The residual sum of squares (RSS) when using the maximum likelihood estimators is

$$\begin{aligned}\text{RSS} &= \sum_{i=1}^n |y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)|^2, \\ &= \sum_{i=1}^n |y_i - \hat{y}_i|^2\end{aligned}$$

- ▶ The RSS is a measure of the model's lack of fit.
- ▶ The model's log likelihood and its RSS are related as follows:

$$\log \mathcal{L} = -\frac{n}{2} (\log(2\pi) - \log(n) + \log(\text{RSS}) + 1)$$

Residual sum of squares and log likelihood

```
rss <- sum(residuals(M1)^2)
n <- length(y)

-(n/2) * (log(2*pi) - log(n) + log(rss) + 1)
#> [1] -206.5784
logLik(M1)
#> 'log Lik.' -206.5784 (df=3)
```

Root mean square error

- ▶ The larger the sample size, the larger the RSS.
- ▶ An alternative to RSS as a measure of model fit is the square root of the mean of the squared residuals, known as the *root mean square error* (RMSE):

$$\text{RMSE} = \sqrt{\frac{\text{RSS}}{n}},$$

- ▶ This is $\hat{\sigma}_{\text{mle}}$.

Mean absolute error

- Related to RMSE is the mean absolute error (MAE), which is the mean of the absolute values of the residuals.

$$\text{MAE} = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$

- In R

```
mean(abs(residuals(M1)))  
#> [1] 11.58012
```


Deviance

- ▶ Deviance is used as a measure of model fit in generalized linear models.
- ▶ Strictly speaking, the deviance of model M_0 is

$$2 (\log \mathcal{L}_s - \log \mathcal{L}_0),$$

where $\log \mathcal{L}_0$ is the log likelihood (at its maximum) of model M_0 , and $\log \mathcal{L}_s$ is a *saturated* model, i.e. one with as many parameters as there are data points.

- ▶ When comparing two models, M_0 and M_1 , the saturated model is the same, and so the difference of the deviances of M_0 and M_1 is

$$\begin{aligned} & (-2 \log \mathcal{L}_0) - (-2 \log \mathcal{L}_1), \\ & \mathcal{D}_0 - \mathcal{D}_1, \end{aligned}$$

and so the deviance of M_0 is usually defined simply as

$$-2 \log \mathcal{L}_0.$$

Differences of deviances

- Differences of deviances are equivalent to log likelihood ratios:

$$\begin{aligned}\mathcal{D}_0 - \mathcal{D}_1 &= -2 \log \mathcal{L}_0 - (-2 \log \mathcal{L}_1), \\ &= -2 (\log \mathcal{L}_0 - \log \mathcal{L}_1), \\ &= -2 \log \left(\frac{\mathcal{L}_0}{\mathcal{L}_1} \right), \\ &= 2 \log \left(\frac{\mathcal{L}_1}{\mathcal{L}_0} \right).\end{aligned}$$

- Clearly, $\frac{\mathcal{L}_1}{\mathcal{L}_0}$ the factor by which the likelihood of model M_1 is greater than that of model M_0 .
- Therefore, the difference of the deviance of models M_0 and M_1 ($\mathcal{D}_0 - \mathcal{D}_1$), gives the (two times) the logarithm of the factor by the likelihood of model M_1 is greater than that of model M_0 .
- The larger $\mathcal{D}_0 - \mathcal{D}_1$, the greater the likelihood of M_1 compared to M_0 .

Logistic regression example

```
cars_df <- mutate(cars, z = dist > median(dist))
M2 <- glm(z ~ speed,
          data = cars_df,
          family = binomial(link = 'logit')
)

logLik(M2)
#> 'log Lik.' -17.73468 (df=2)
deviance(M2)
#> [1] 35.46936
logLik(M2) * -2
#> 'log Lik.' 35.46936 (df=2)
```

Conditional probability in logistic regression

- The model in a logistic regression (with one predictor) is

$$y_i \sim \text{Bernoulli}(\theta_i), \quad \text{for } i \in 1 \dots n$$
$$\log \left(\frac{\theta_i}{1 - \theta_i} \right) = \beta_0 + \beta_1 x_i$$

- The conditional probability of $y_1, y_2 \dots y_n$ given $x_1, x_2 \dots x_n$ is

$$\prod_{i=1}^n \theta_i^{y_i} (1 - \theta_i)^{1-y_i},$$

where each θ_i is

$$\log \left(\frac{\theta_i}{1 - \theta_i} \right) = \beta_0 + \beta_1 x_i$$

Conditional probability in logistic regression

- The logarithm of the conditional probability of $y_1, y_2 \dots y_n$ is

$$\begin{aligned} \log \left(\prod_{i=1}^n \theta_i^{y_i} (1 - \theta_i)^{1-y_i} \right) &= \sum_{i=1}^n \log \left(\theta_i^{y_i} (1 - \theta_i)^{1-y_i} \right), \\ &= \sum_{i=1}^n (y_i \log \theta_i + (1 - y_i) \log(1 - \theta_i)), \\ &= \sum_{i=1}^n y_i \log \theta_i + \sum_{i=1}^n (1 - y_i) \log(1 - \theta_i) \end{aligned}$$

Conditional probability in logistic regression

```
theta <- predict(M2, type = 'response')
sum(log(theta[cars_df$z])) + sum(log(1-theta[!cars_df$z]))
#> [1] -17.73468
```

```
z <- pull(cars_df, z)
sum(z * log(theta) + (1-z) * log(1 - theta))
#> [1] -17.73468
```

Deviance residuals

- ▶ Deviance residuals are values such that their sum of squares is equal to the model's deviance.
- ▶ We know that the sum, for $i \in 1 \dots n$, of the following is the log likelihood:

$$y_i \log \theta_i + (1 - y_i) \log(1 - \theta_i),$$

and so the sum of the following, for $i \in 1 \dots n$, is the deviance:

$$-2 (y_i \log \theta_i + (1 - y_i) \log(1 - \theta_i)).$$

- ▶ So the sum of the *squares* of the following, for $i \in 1 \dots n$, is the deviance:

$$\sqrt{-2 (y_i \log \theta_i + (1 - y_i) \log(1 - \theta_i))}.$$

- ▶ All of these values will necessarily be positive.
- ▶ It is conventional for deviance residuals to be negative when $y_i = 0$ and positive when $y_i = 1$.

Deviance residuals

```
d <- sqrt( -2 * (z * log(theta) + (1-z) * log(1 - theta)))  
sum(d^2)  
#> [1] 35.46936
```

```
d[c(1, 25, 35, 50)]  
#>           1           25           35           50  
#> 0.05724272 1.00995907 0.71599367 0.11291237
```

```
residuals(M2)[c(1, 25, 35, 50)]  
#>           1           25           35           50  
#> -0.05724272 -1.00995907 0.71599367 0.11291237  
z[c(1, 25, 35, 50)]  
#> [1] FALSE FALSE  TRUE  TRUE  
(ifelse(z, 1, -1) * d)[c(1, 25, 35, 50)]  
#>           1           25           35           50  
#> -0.05724272 -1.00995907 0.71599367 0.11291237
```